

Deep Learning for NLP

- Oliver Guhr
I833 SS2019

How is the pace of the lectures so far?

too slow / too fast / just right

a brief recap of the last lecture

Word Representations

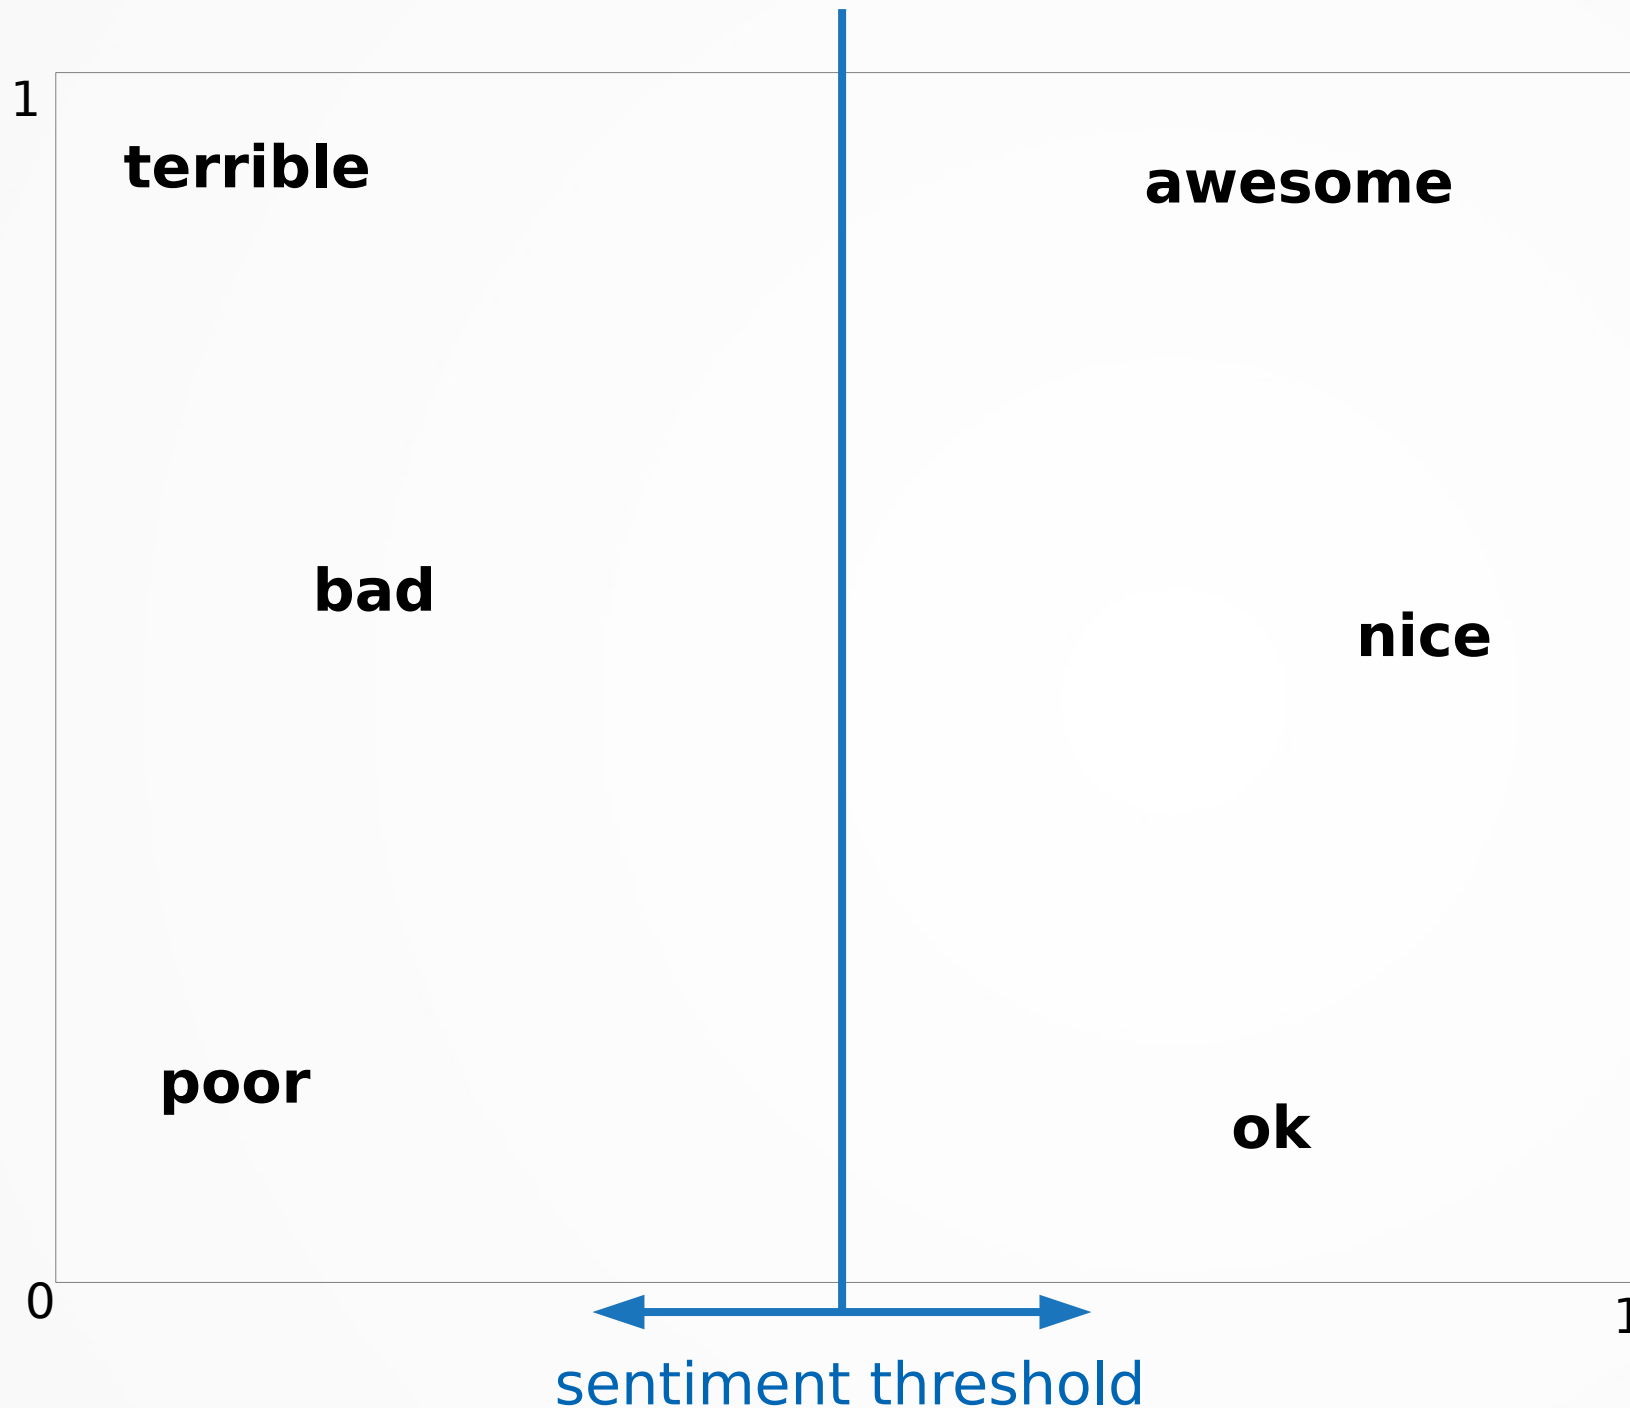
Encoding

Instead of encoding single characters

h	0	0	0	1
e	0	0	1	0
l	0	1	0	0
o	1	0	0	0

You can also encode words, this is also called „Bag-Of-Words (BOW)“

hello	0	0	0	1
my	0	0	1	0
name	0	1	0	0
is	1	0	0	0



Now a network can distinguish between positive and negative words by learning a threshold.

$$v_{ok} = [0.75, 0.15]$$

$$v_{nice} = [0.85, 0.50]$$

$$v_{poor} = [0.15, 0.18]$$

$$v_{terrible} = [0.10, 0.91]$$

Distributional Hypothesis

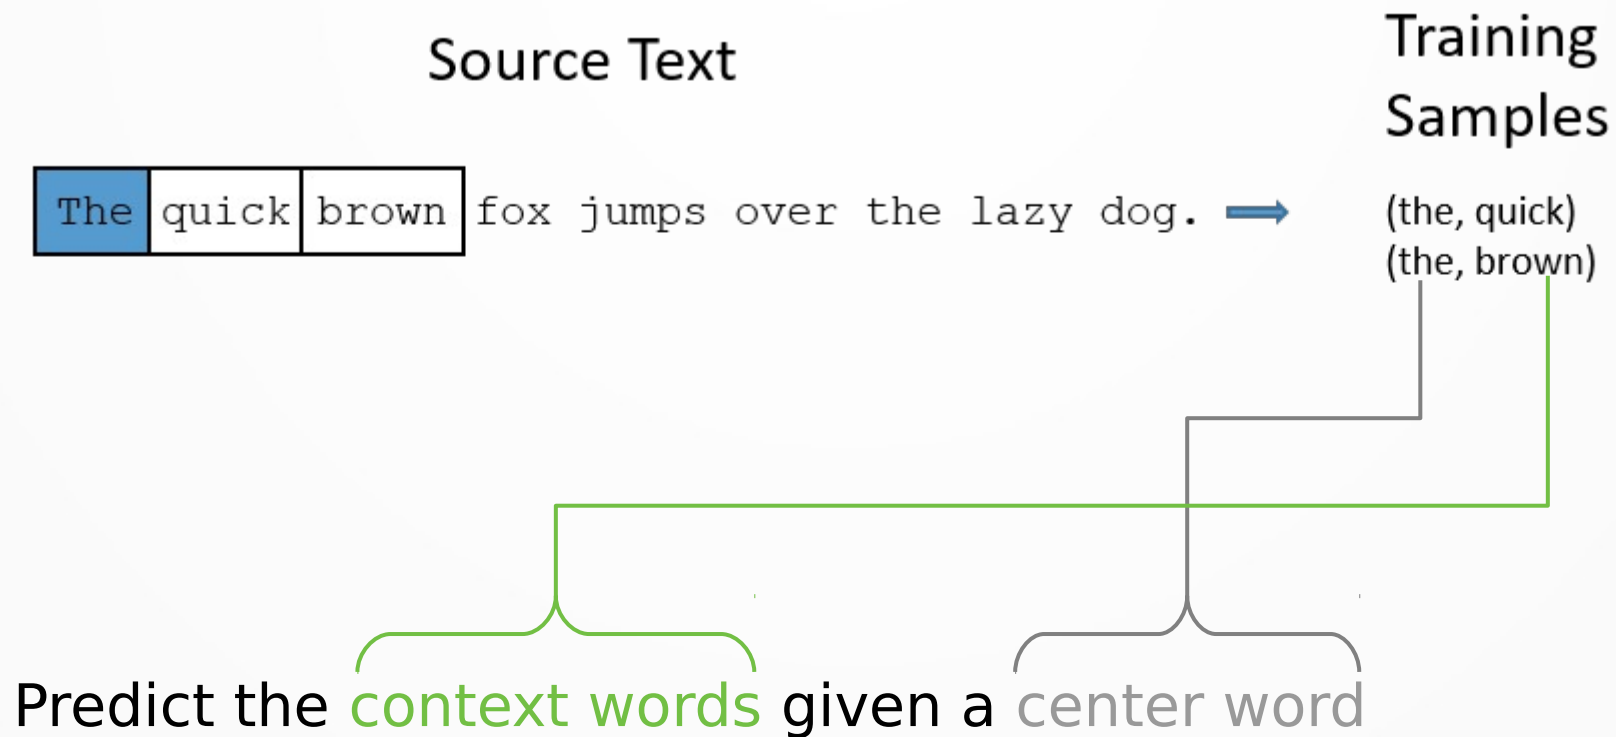
Words that occur in the same contexts tend to have similar meanings.

Harris (1954)

A word is characterized by the company it keeps.

Firth (1957)

Skip-Gram



Supervised and unsupervised learning

- **Supervised Learning** uses a set of labeled training examples.
 - List of e-mails that are labeled as spam / not spam
- **Unsupervised Learning** generates training examples from a plain -unlabeled- text corpus, so the training becomes a supervised problem.
 - Predict context words (Skip Gram)
 - Predict center words (Cbow)
 - Predict masked words (Bert)

Preprocessing

- Format your text to be predictable and analyzable
- It often has a significant impact on the performance
- Depending on the domain and your model different steps may be required
- For example:
 - Cleaning not useful characters and words
 - Transform words into a standardized form
 - Clipping your data to equal length

Scores: Accuracy

- Accuracy =
$$\frac{tp + tn}{tp + tn + fp + fn}$$
- Accuracy =
$$\frac{\text{number correctly predicted samples}}{\text{total number of samples}}$$

Scores: F1

$$\textit{precision} = \frac{\textit{true positives}}{\textit{true positives} + \textit{false positives}}$$

$$\textit{recall} = \frac{\textit{true positives}}{\textit{true positives} + \textit{false negatives}}$$

$$F_1 = 2 * \frac{\textit{precision} * \textit{recall}}{\textit{precision} + \textit{recall}}$$

How to optimize:

- 1) create a train/test split
- 2) Train your model (start with a simple model!)
- 3) measure its performance
- 4) optimize your model
- 5) Go to 2 :)

Identify offensive language

using word vectors and
FastText




Goal for Today


- Look we look at deep language models
 - How do they perform?
 - How do they work?
 - Some of the basic ideas behind those models.
 - How to use these models.
 - Applications :)

Deep Language Models

GLUE Benchmark Results



SQUAD 2.0



Deep Language Models

- In 2018 several Ideas led to new models
 - [Semi-supervised Sequence Learning](#) Andrew Dai, Quoc Le
 - [ELMo](#) Peters et al.
 - [ULMFiT](#) Howard, Ruder
 - [OpenAI Transformer](#) Radford, Narasimhan, Salimans, Sutskever
 - [Transformer](#) Vaswani et al.

Deep Language Models

- Google's BERT (October 2018)
 - BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding
- Google/CMU's Transformer-XL (January 2019)
 - Transformer-XL: Attentive Language Models Beyond a Fixed-Length Context
- OpenAI's GPT (June, 2018)
 - Improving Language Understanding by Generative Pre-Training
- OpenAI's GPT-2 (February, 2019)
 - Language Models are Unsupervised Multitask Learners

Model Parameters

Model	Parameters
Medium LSTM	10 Million
ELMo	90 Million
GTP	110 Million
Bert Base	110 Million
Bert Large	340 Million
GTP-2	1500 Million

BERT

Bidirectional Encoder Representations from Transformers



Bert

- What can you do with Bert?
- Some applications:
 - Named Entity Recognition
 - Text Classification
 - Fact Checking
 - Text Summarization
 - Text Generation
 - Question Answering (Full Text and Multiple Choice)
 - Translation

Bert

- Training Process
 - Pre-train a model on plain text
 - Choose a task specific labeled data set
 - Retrain the model with this data set
- Use the same pre trained model for all tasks
 - Classification
 - Named Entity Recognition
 - Question Answering etc.

1 - Semi-supervised training on large amounts of text (books, wikipedia..etc).

The model is trained on a certain task that enables it to grasp patterns in language. By the end of the training process, BERT has language-processing abilities capable of empowering many models we later need to build and train in a supervised way.

Semi-supervised Learning Step

Model:



Dataset:



Objective:

Predict the masked word
(language modeling)

1 - **Semi-supervised** training on large amounts of text (books, wikipedia..etc).

The model is trained on a certain task that enables it to grasp patterns in language. By the end of the training process, BERT has language-processing abilities capable of empowering many models we later need to build and train in a supervised way.

Semi-supervised Learning Step

Model:



Dataset:



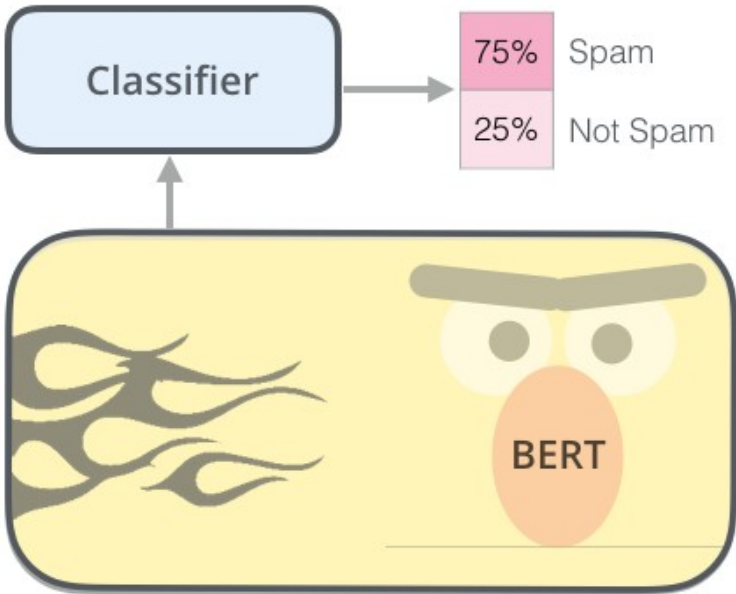
Objective:

Predict the masked word
(language modeling)

2 - **Supervised** training on a specific task with a labeled dataset.

Supervised Learning Step


Model:
(pre-trained
in step #1)




Dataset:

Email message	Class
Buy these pills	Spam
Win cash prizes	Spam
Dear Mr. Atreides, please find attached...	Not Spam

Two pre-trained sizes



BERT_{BASE}



BERT_{LARGE}

Pre Trained Bert

- English
 - BERT-Large, Uncased (Whole Word Masking): 24-layer, 1024-hidden, 16-heads, 340M parameters
 - BERT-Large, Cased (Whole Word Masking): 24-layer, 1024-hidden, 16-heads, 340M parameters
 - BERT-Base, Uncased: 12-layer, 768-hidden, 12-heads, 110M parameters
 - BERT-Large, Uncased: 24-layer, 1024-hidden, 16-heads, 340M parameters
 - BERT-Base, Cased: 12-layer, 768-hidden, 12-heads , 110M parameters
 - BERT-Large, Cased: 24-layer, 1024-hidden, 16-heads, 340M parameters
- Multi Language
 - BERT-Base, Multilingual Cased (New, recommended): 104 languages, 12-layer, 768-hidden, 12-heads, 110M parameters
 - BERT-Base, Multilingual Uncased (Orig, not recommended) (Not recommended, use Multilingual Cased instead): 102 languages, 12-layer, 768-hidden, 12-heads, 110M parameters
- Chinese
 - BERT-Base, Chinese: Chinese Simplified and Traditional, 12-layer, 768-hidden, 12-heads, 110M parameters

Pre Trained Bert

Explained on the next slides

- English
 - BERT-Large, Uncased (Whole Word Masking): 24-layer, 1024-hidden, 16-heads, 340M parameters
 - BERT-Large, Cased (Whole Word Masking): 24-layer, 1024-hidden, 16-heads, 340M parameters
 - BERT-Base, Uncased: 12-layer, 768-hidden, 12-heads, 110M parameters
 - BERT-Large, Uncased: 24-layer, 1024-hidden, 16-heads, 340M parameters
 - BERT-Base, Cased: 12-layer, 768-hidden, 12-heads, 110M parameters
 - BERT-Large, Cased: 24-layer, 1024-hidden, 16-heads, 340M parameters
- Multi Language
 - BERT-Base, Multilingual Cased (New, recommended): 104 languages, 12-layer, 768-hidden, 12-heads, 110M parameters
 - BERT-Base, Multilingual Uncased (Orig, not recommended) (Not recommended, use Multilingual Cased instead): 102 languages, 12-layer, 768-hidden, 12-heads, 110M parameters
- Chinese
 - BERT-Base, Chinese: Chinese Simplified and Traditional, 12-layer, 768-hidden, 12-heads, 110M parameters

Bert

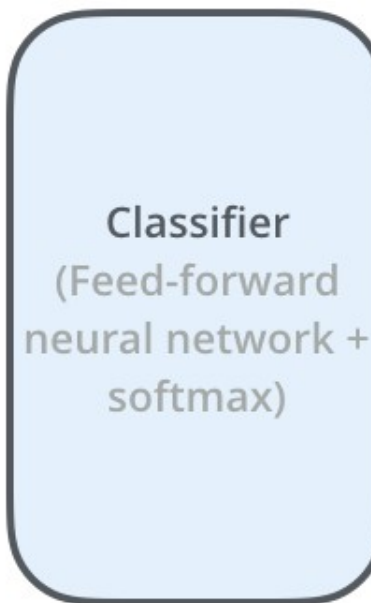
- Issues with non English models
 - One model for 102 languages
 - Pre-trained on Wikipedia content
 - English Models are pre-trained on books and Wikipedia
 - Train you own language model will cost ca. 500\$
 - Collecting this data for German language would be a great task fo a research seminar :)

Text Classification

Input
Features

Output
Prediction


Help Prince Mayuko Transfer
Huge Inheritance



Two sizes

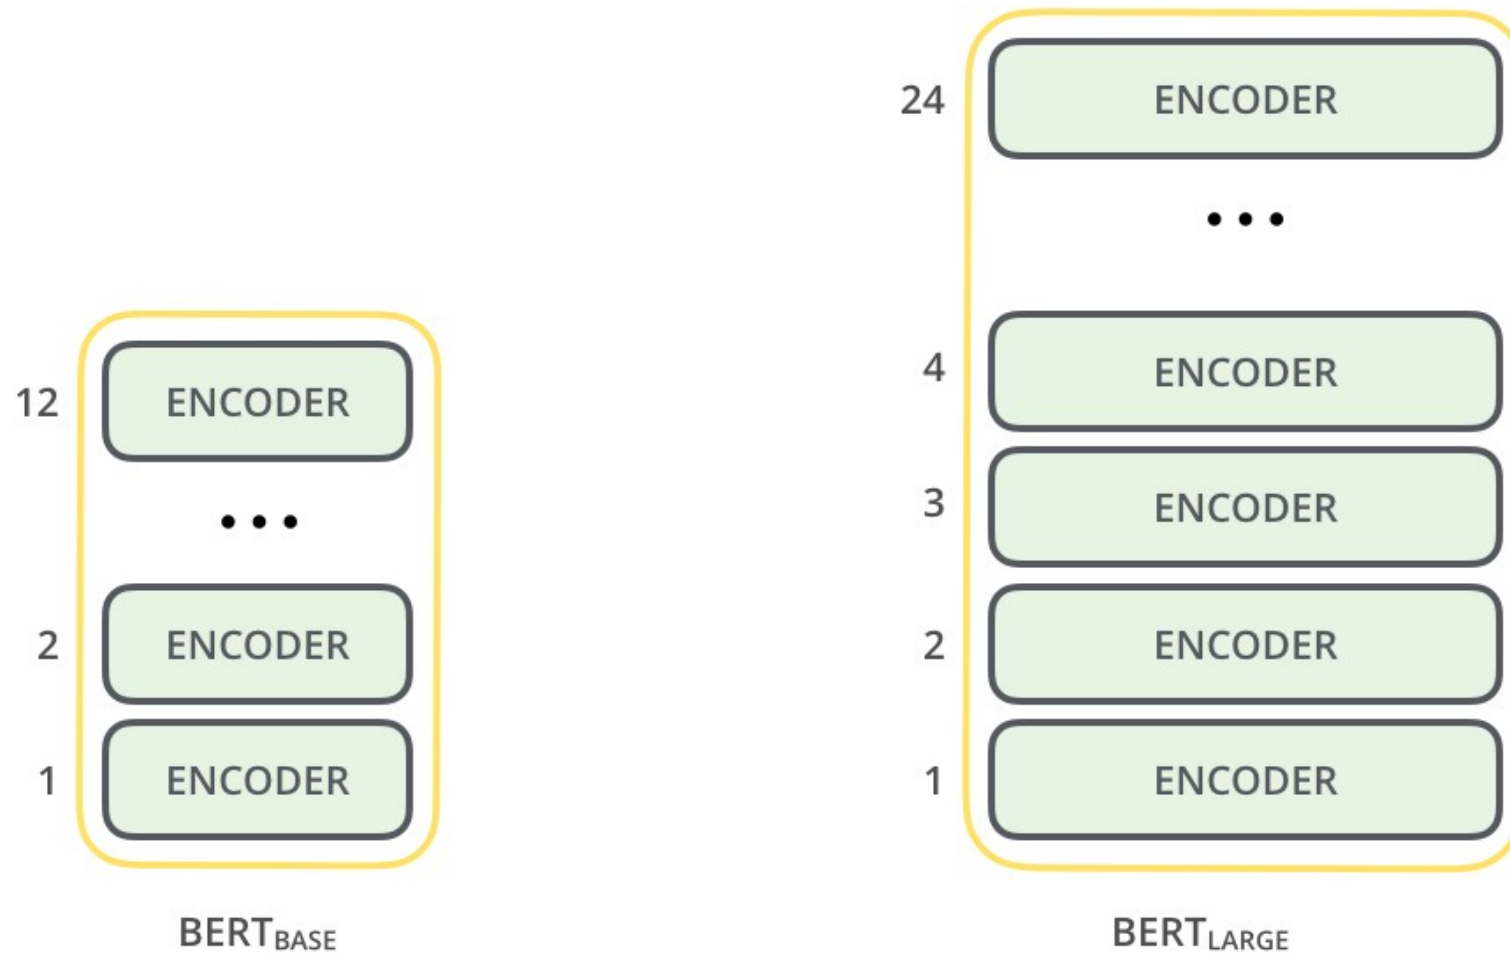


BERT_{BASE}



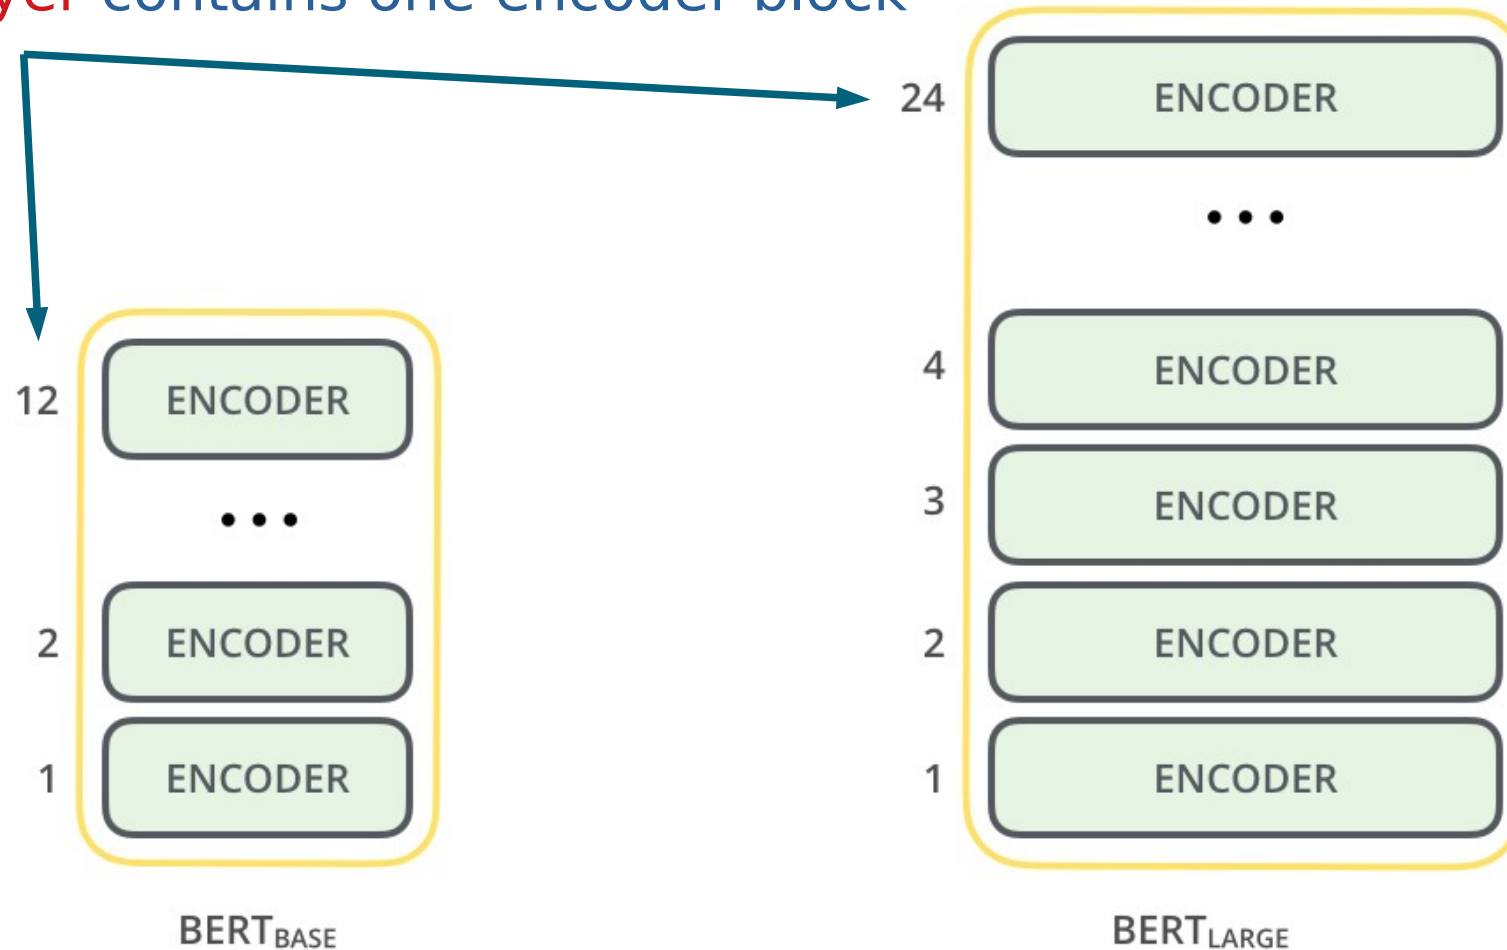
BERT_{LARGE}

Bert Model




Bert Model


Each **layer** contains one encoder block



How to process sequences?

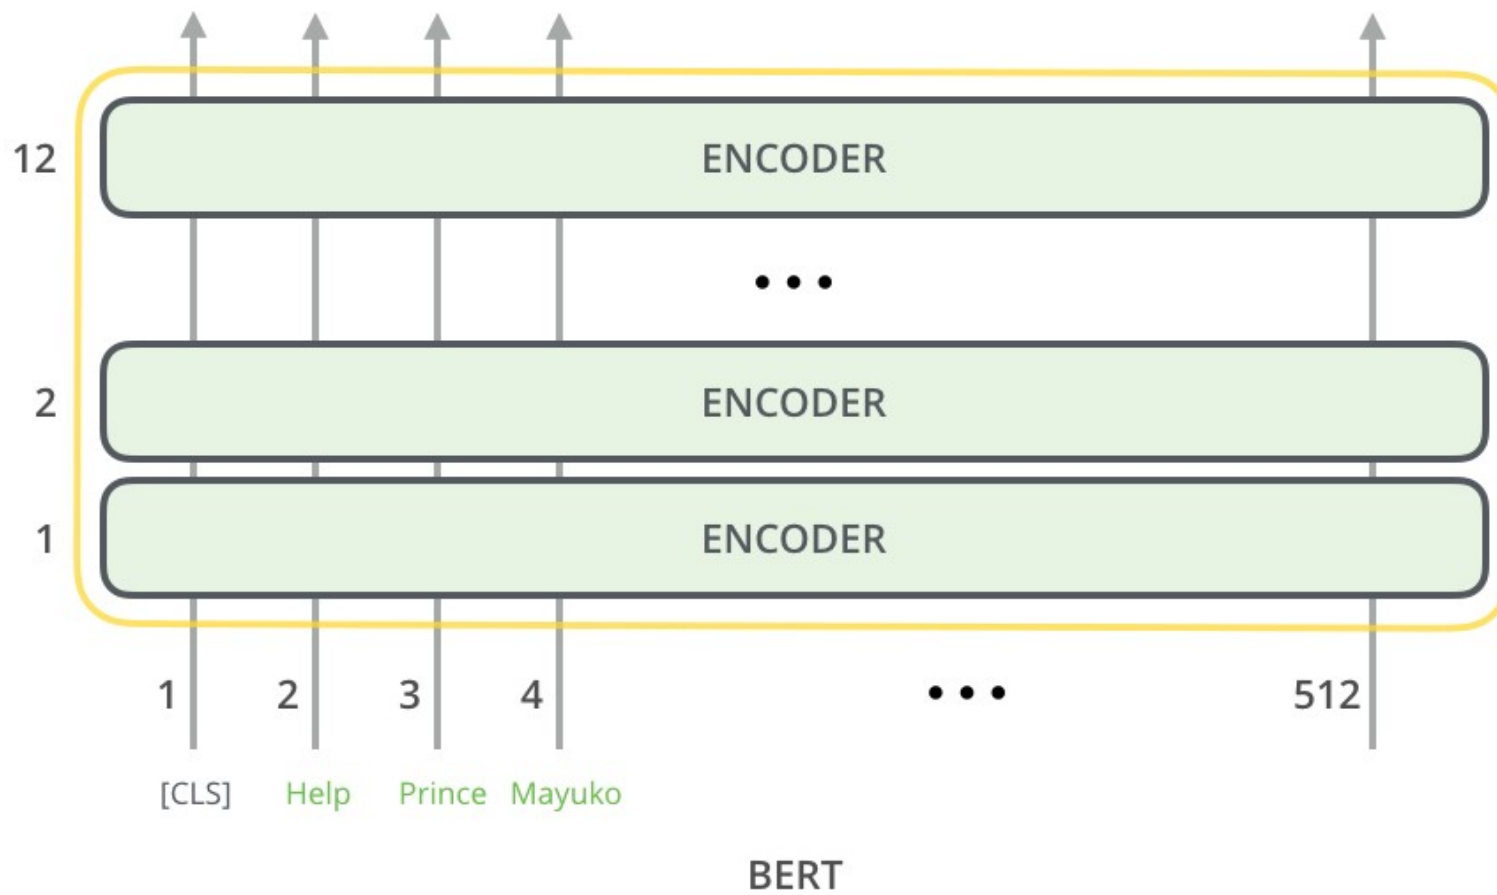


How to process sequences?




This Bert model can process sequences up to 512 tokens.

Bert



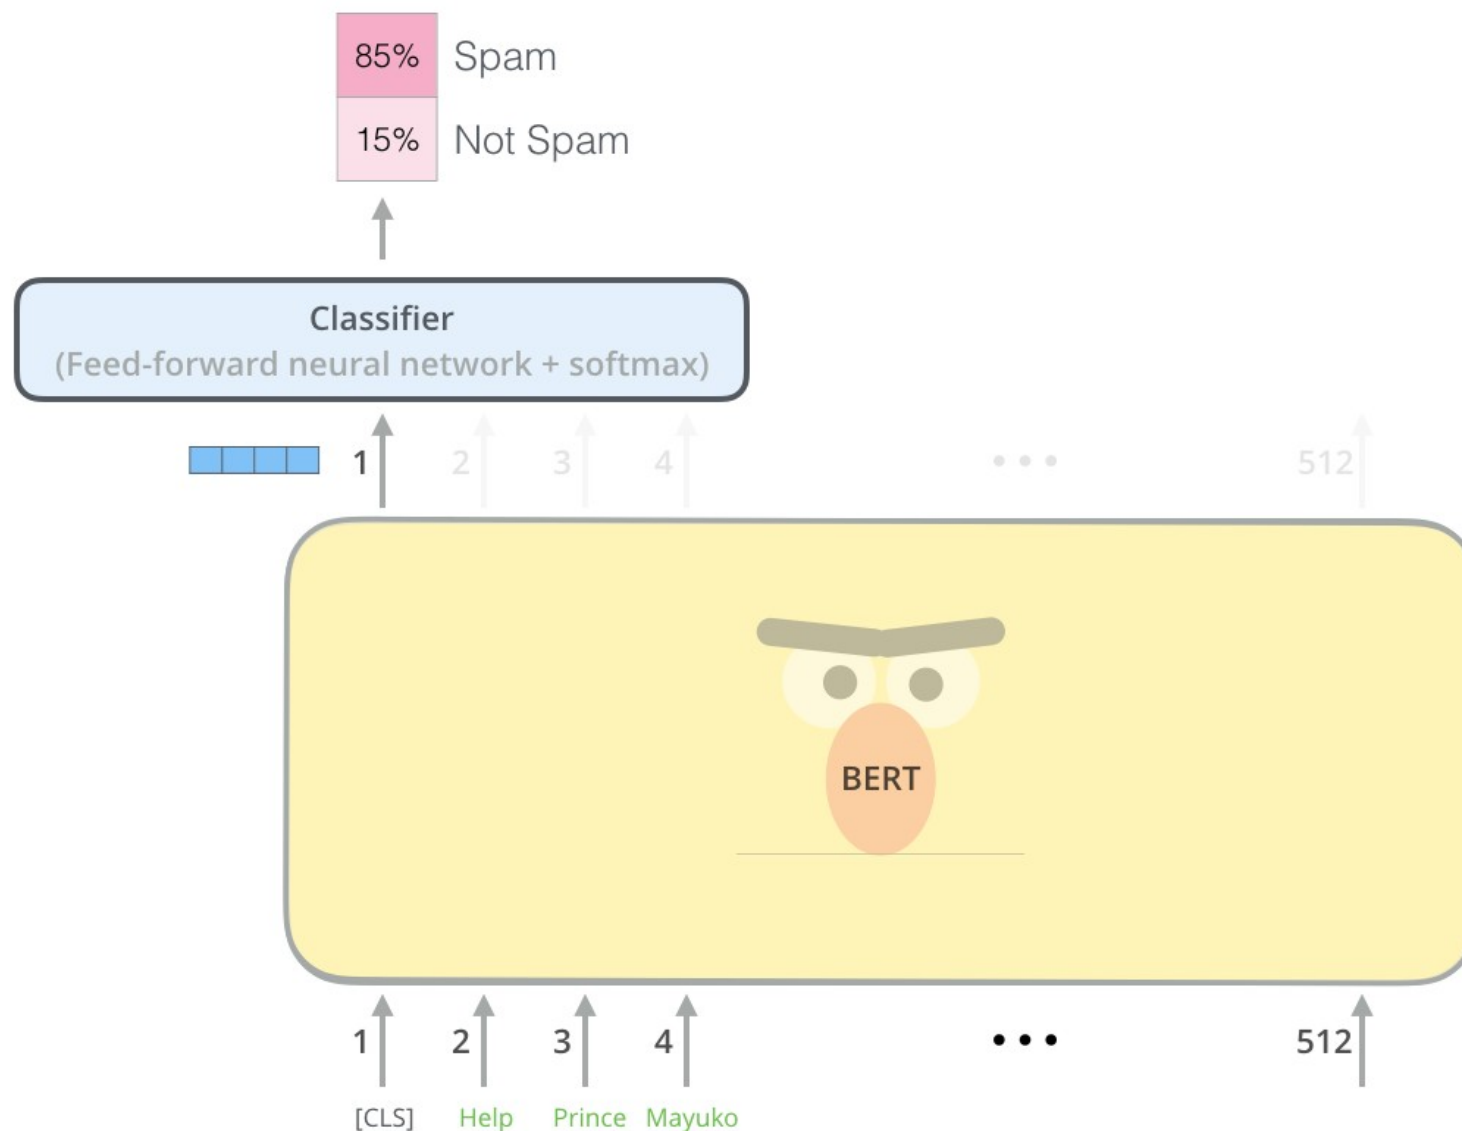
Bert

Each token generates a vector with the length of the **hidden size**.



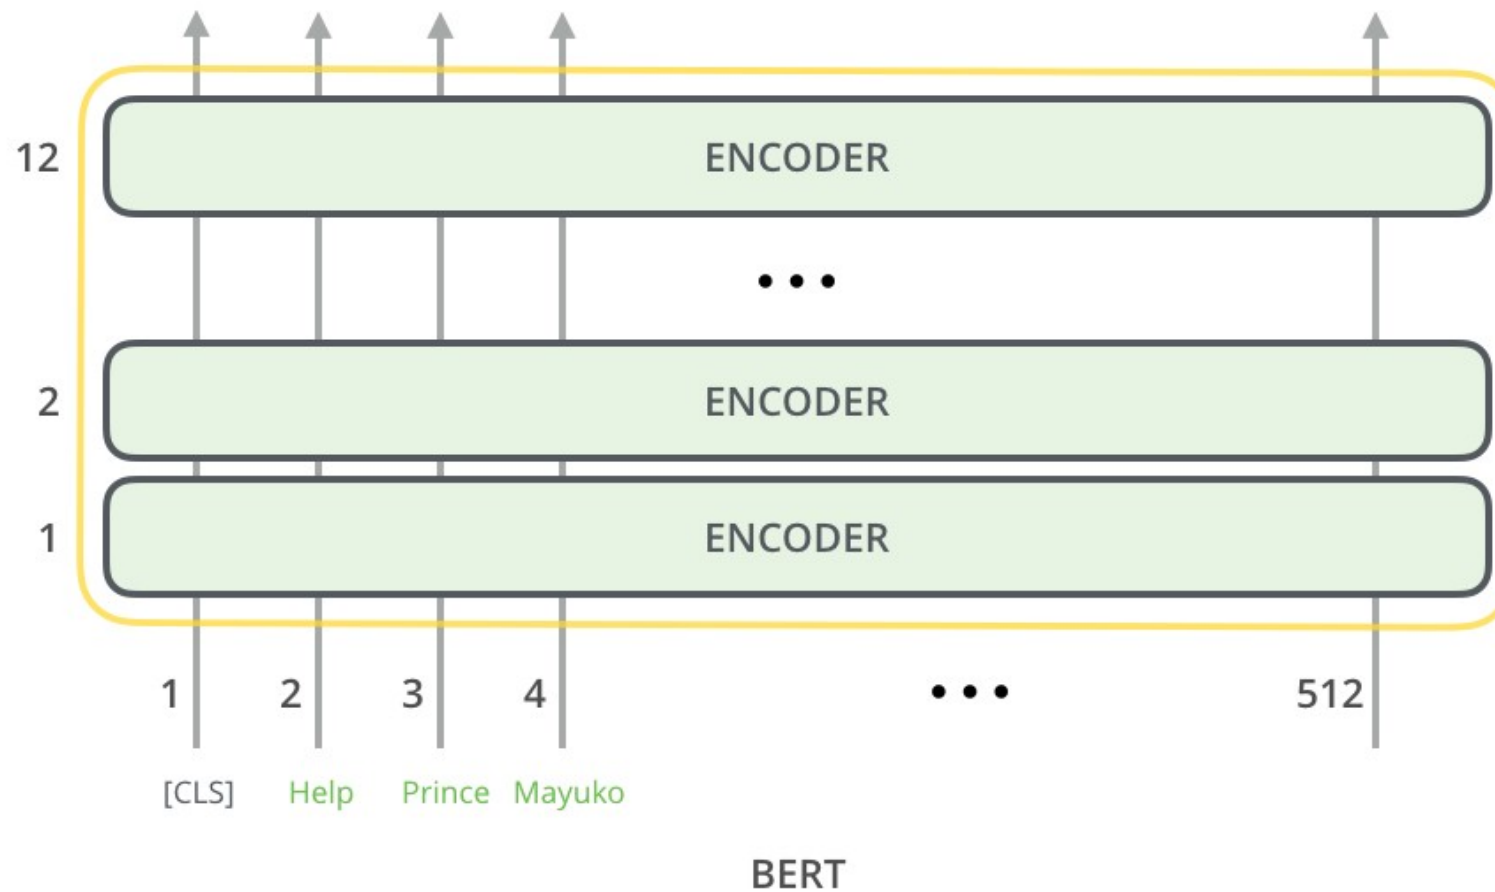
BERT

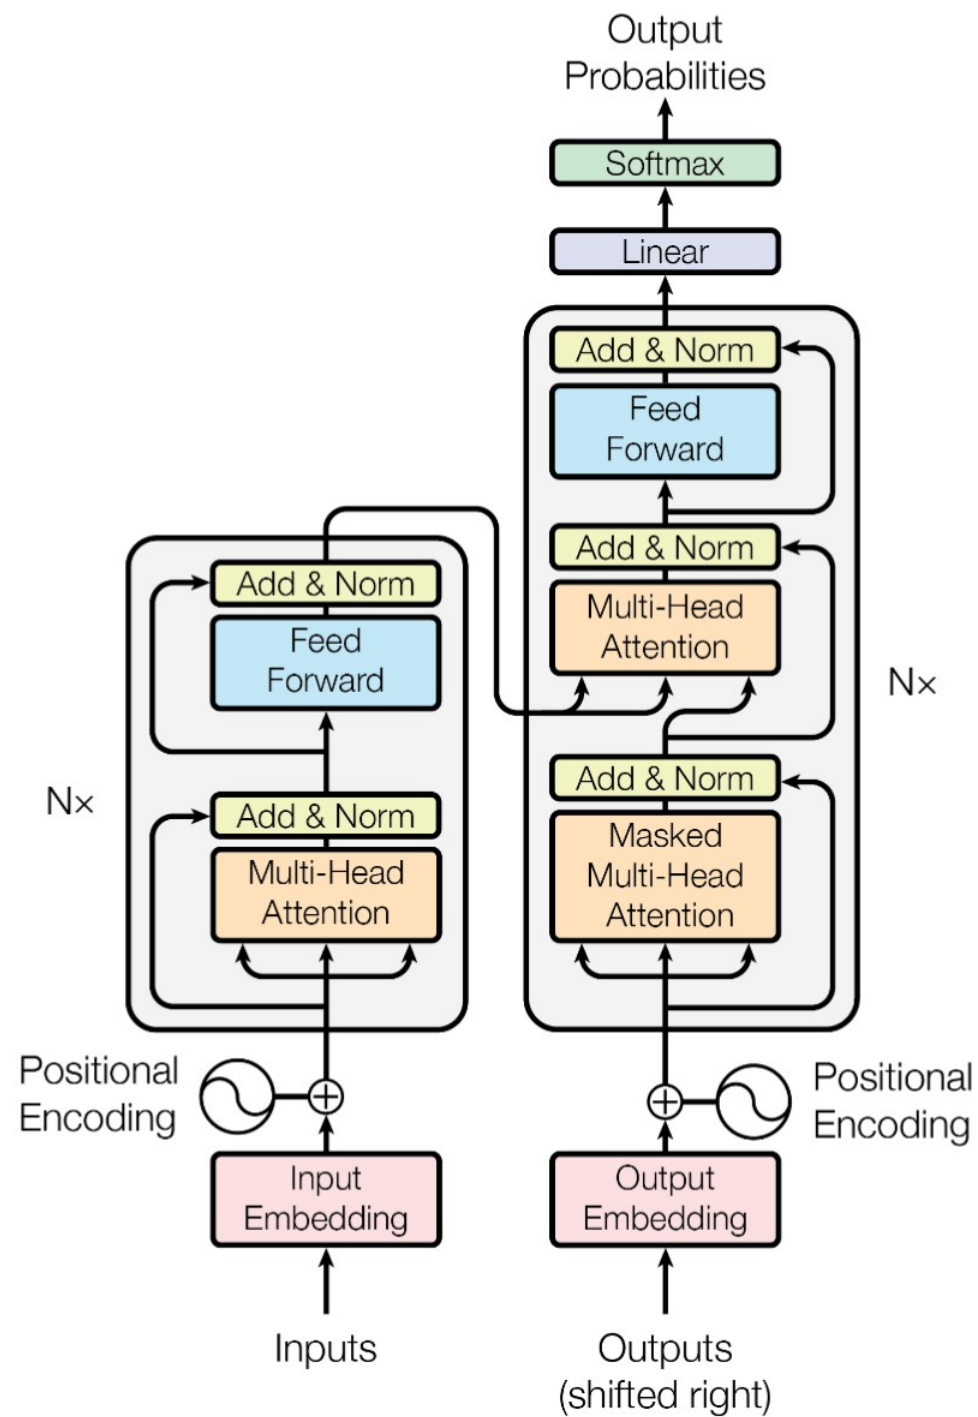
Classification with Bert

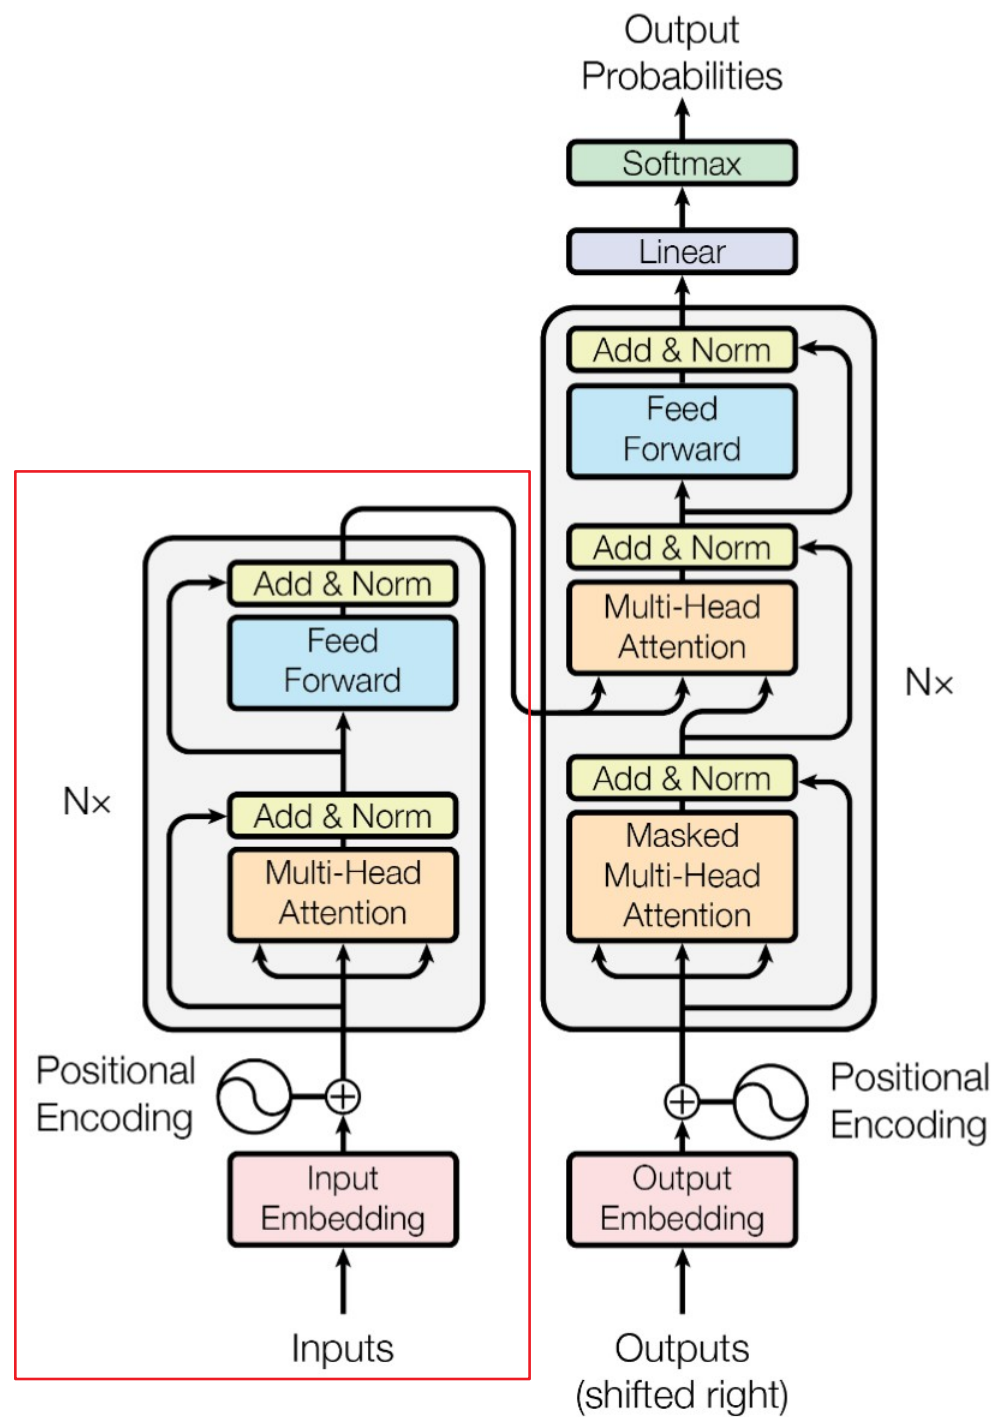


Attention and Transformer

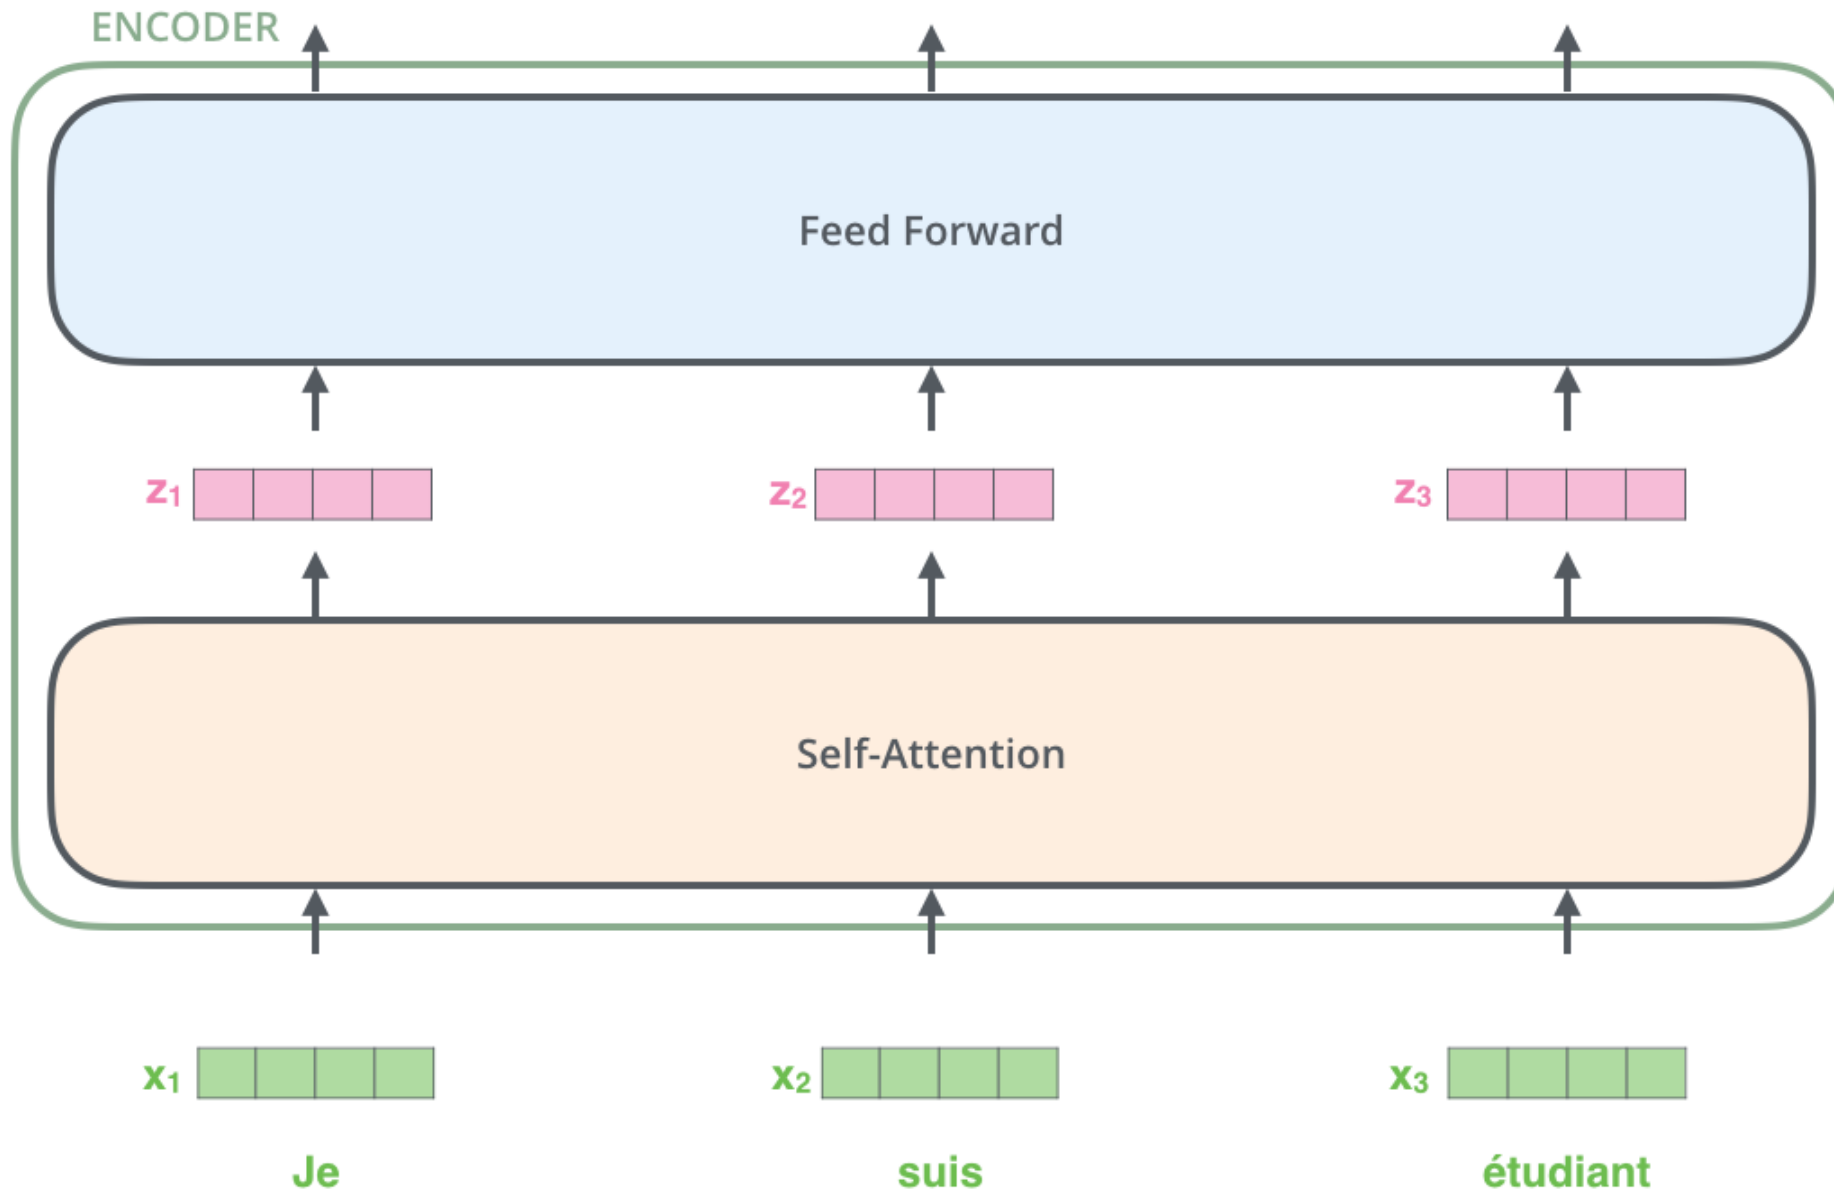
Berts Encoder







Transfomer




Transformer Attention

$$A(\underbrace{Q}_{\text{Query}}, \underbrace{K}_{\text{Key}}, \underbrace{V}_{\text{Value}}) = \text{softmax}(QK^T)V$$

Take the current **word or token**, find the most similar **Key** and return the corresponding **value**.

Transformer: Multi Head Attention



Transformer

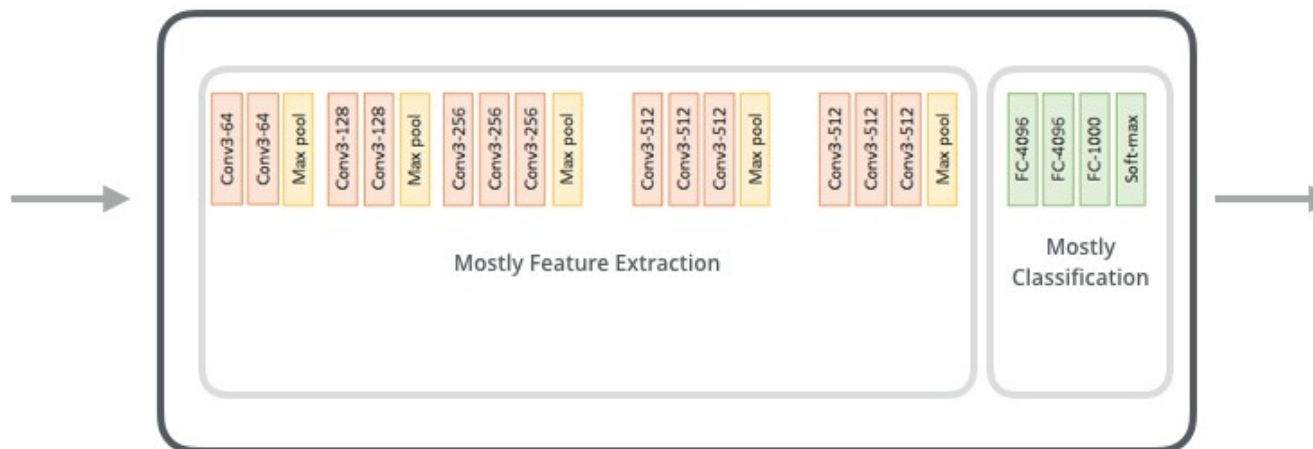
- Paper
 - Attention is all you need. Vaswani et al.
- Good Read
 - Jay Alammars [The Illustrated Transformer](#)
- Conference Talk:
 - Attention is all you need attentional neural network models by Łukasz Kaiser

Similarity to Conv Nets

Input
Features



VGG-16



Output
Prediction

0.2%	Kit fox
0.1%	English setter
95%	Egyptian cat
1%	Great Dane
	...
0%	Hotdog

Training Bert

Pre Training Bert

- pre trained models are also called **language models**
- Compared to FastText Berts language model can distinguish between contexts
 - “river bank” vs “financial bank”
- To create them, Bert used two methods:
 - Task One: Mask Words
 - Task Two: Next Sentence Prediction

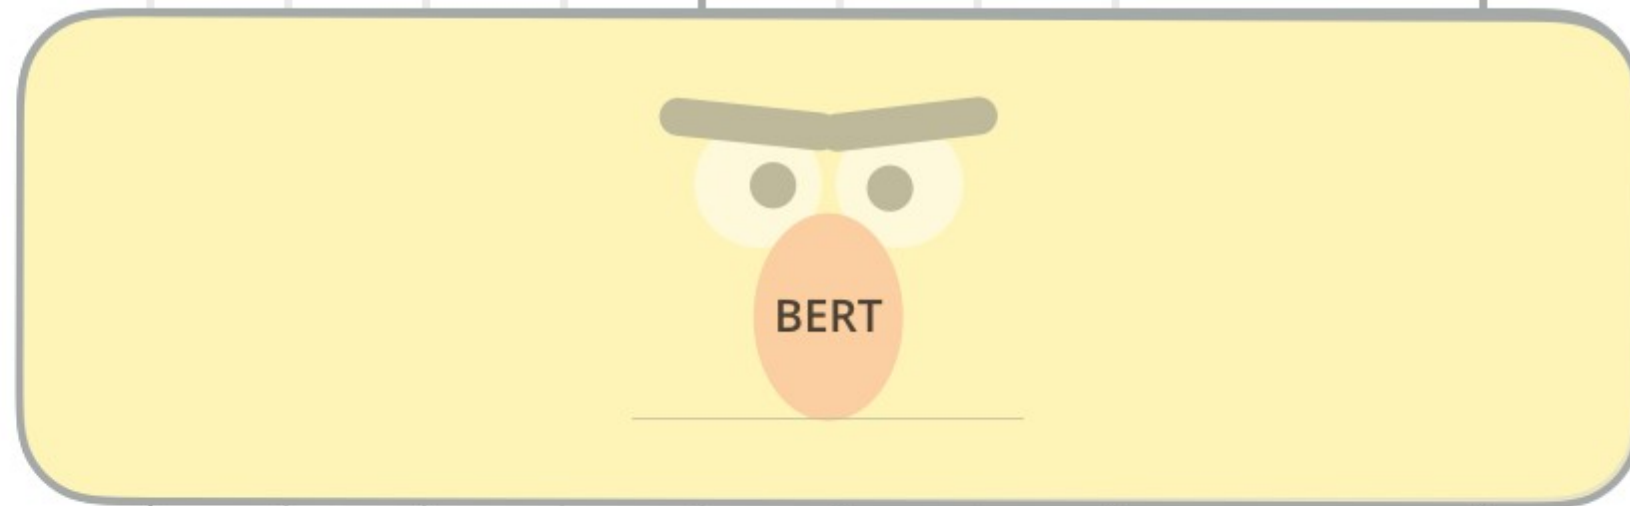
Use the output of the masked word's position to predict the masked word

Possible classes:
All English words

0.1%	Aardvark
...	...
10%	Improvisation
...	...
0%	Zyzzzyva

FFNN + Softmax

1 2 3 4 5 6 7 8 ... 512



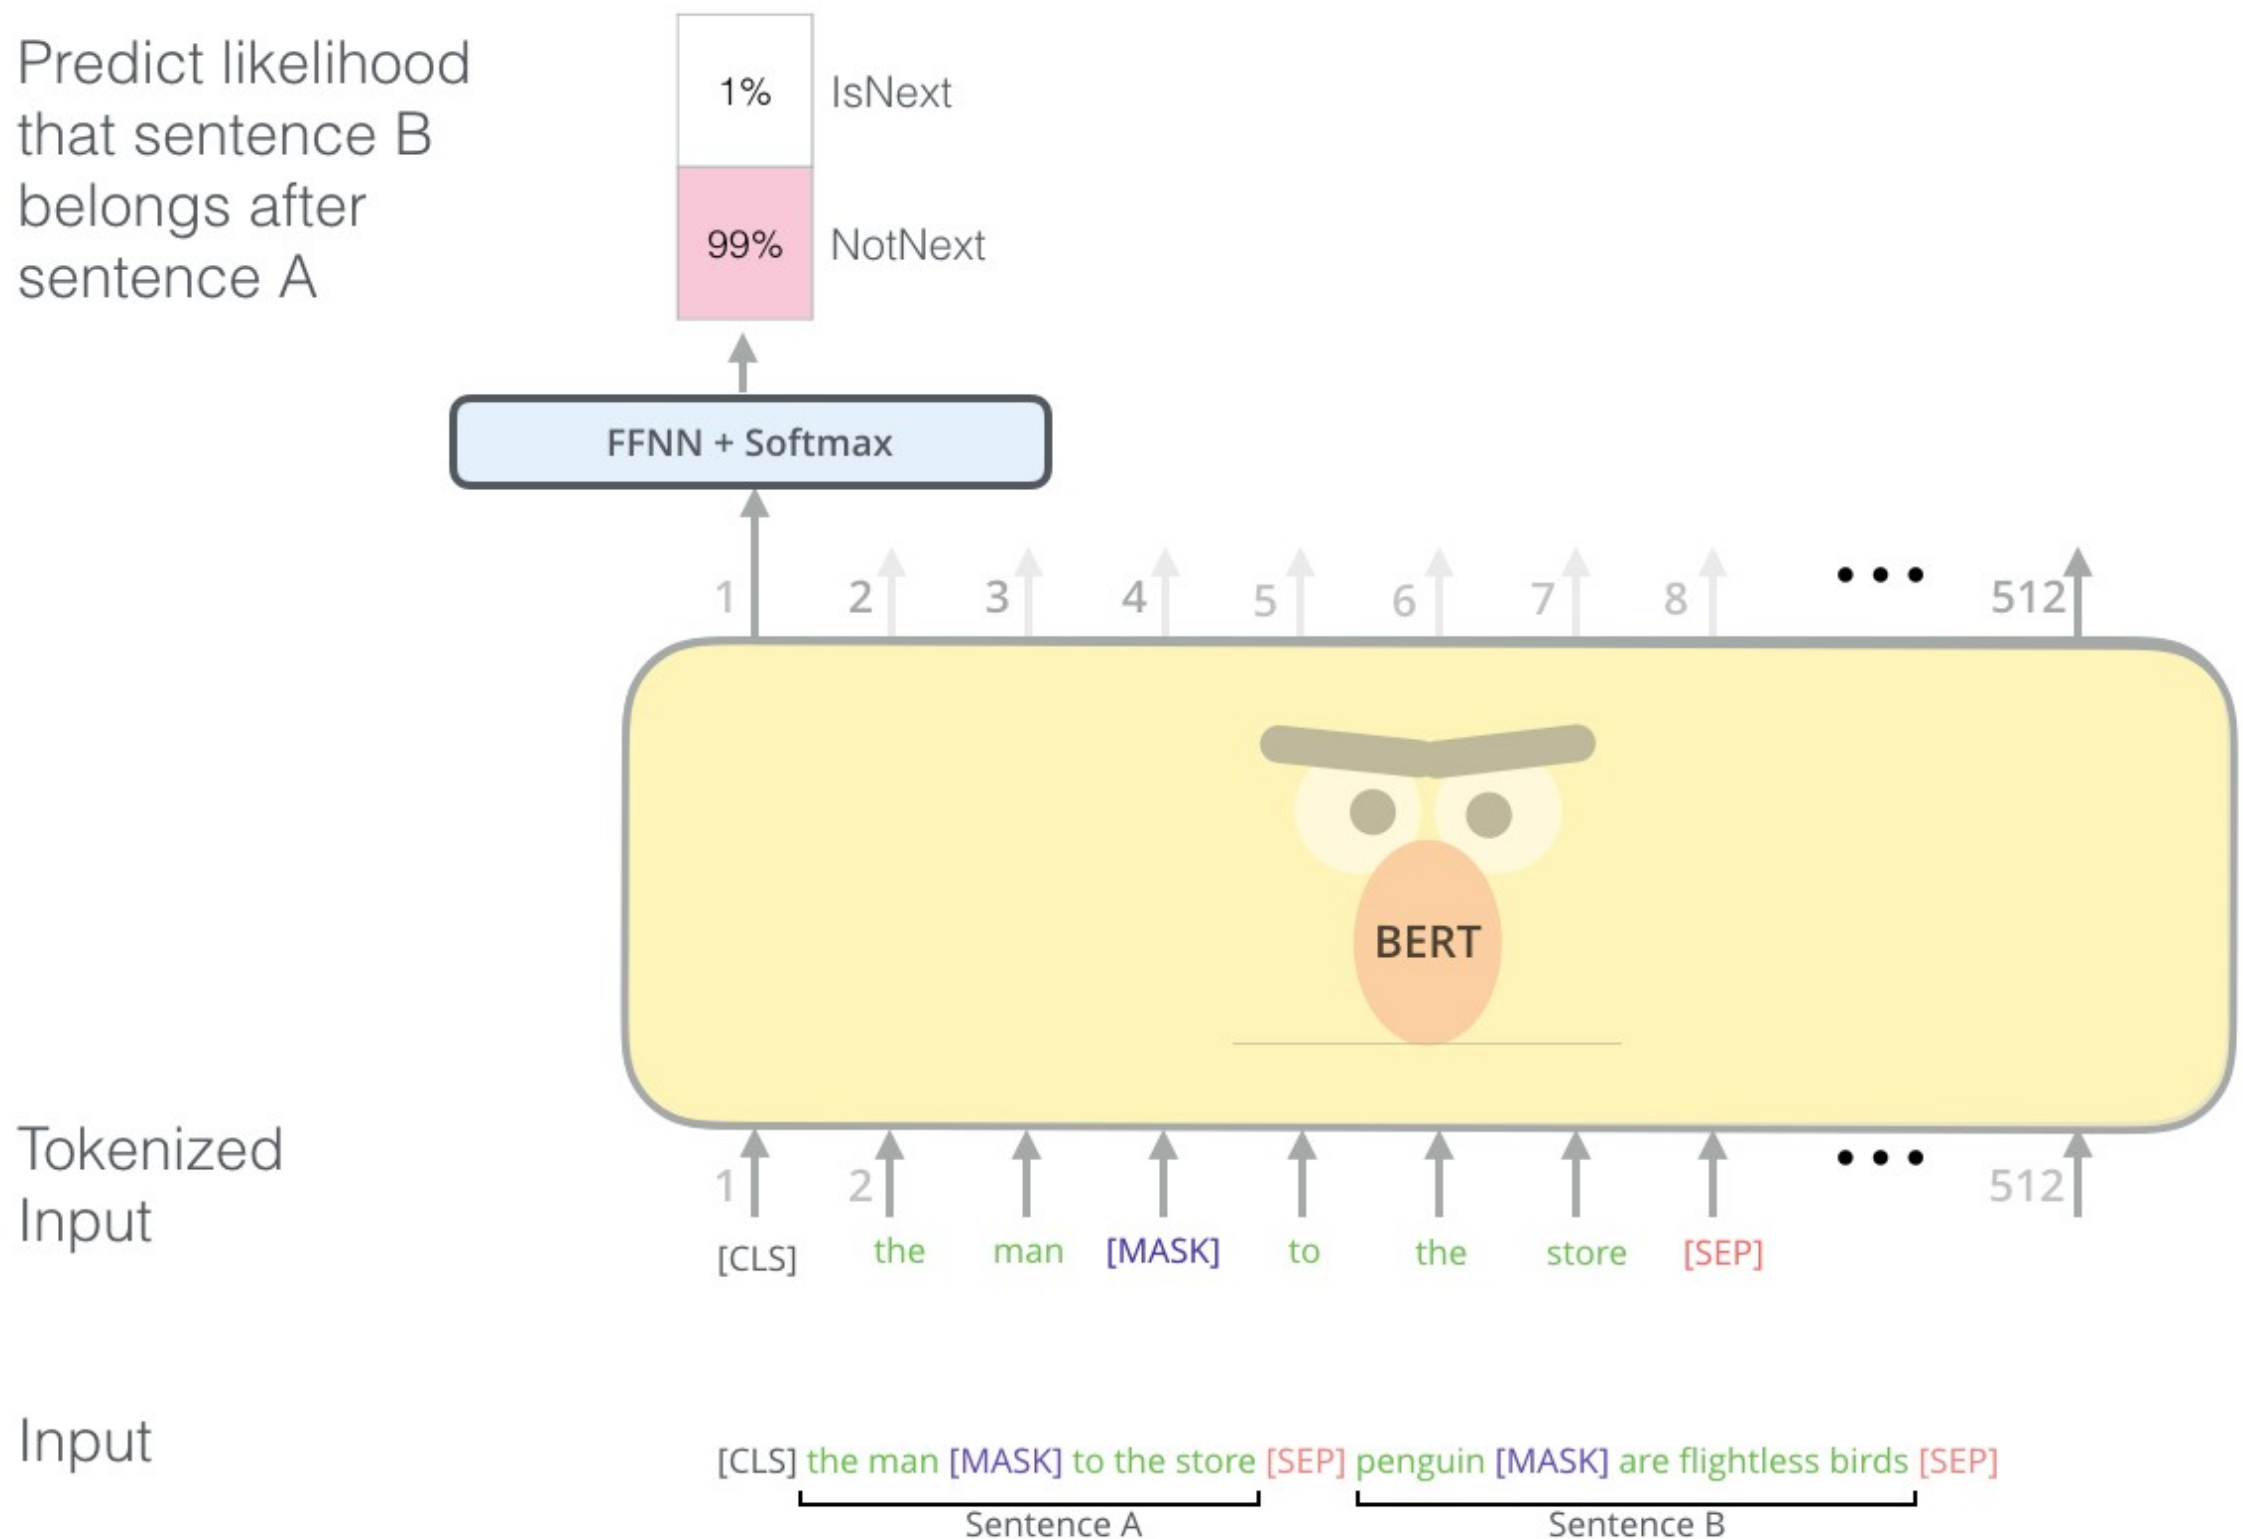
Randomly mask
15% of tokens

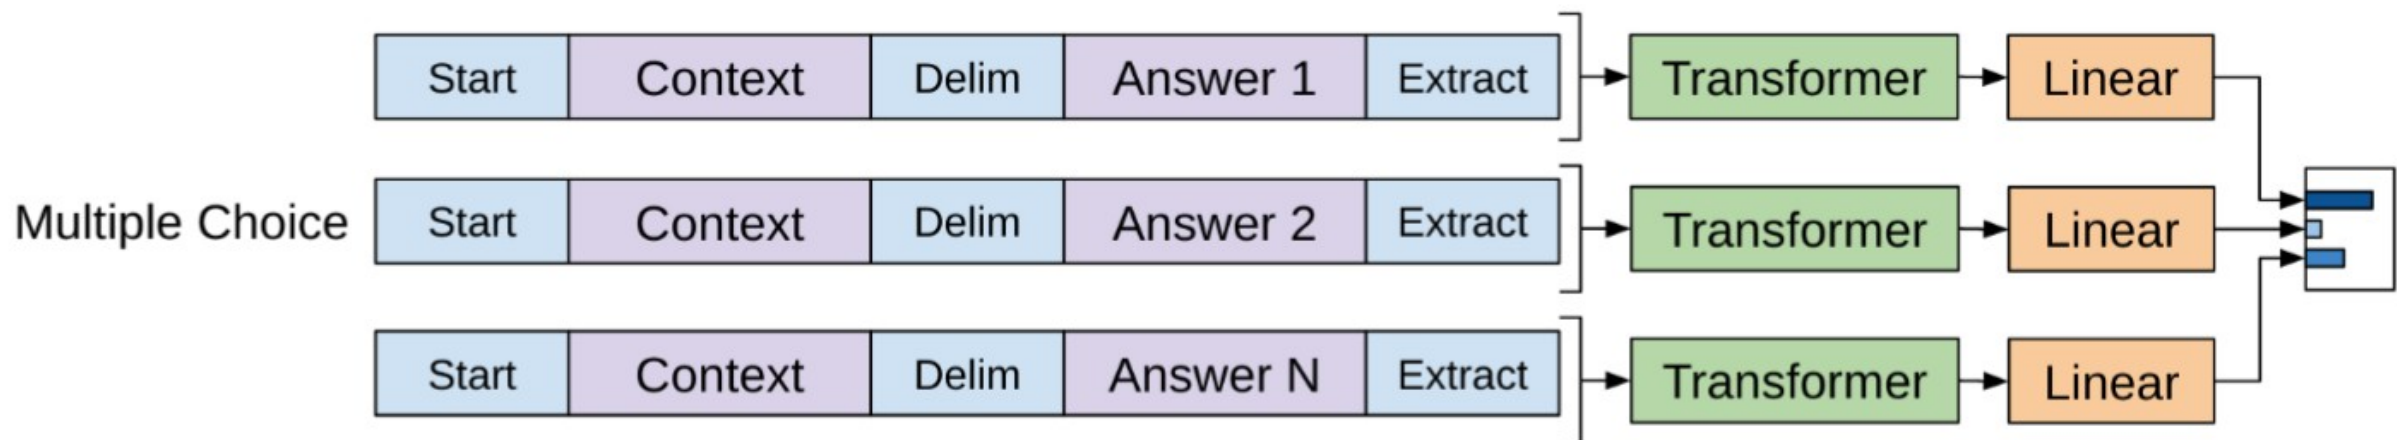
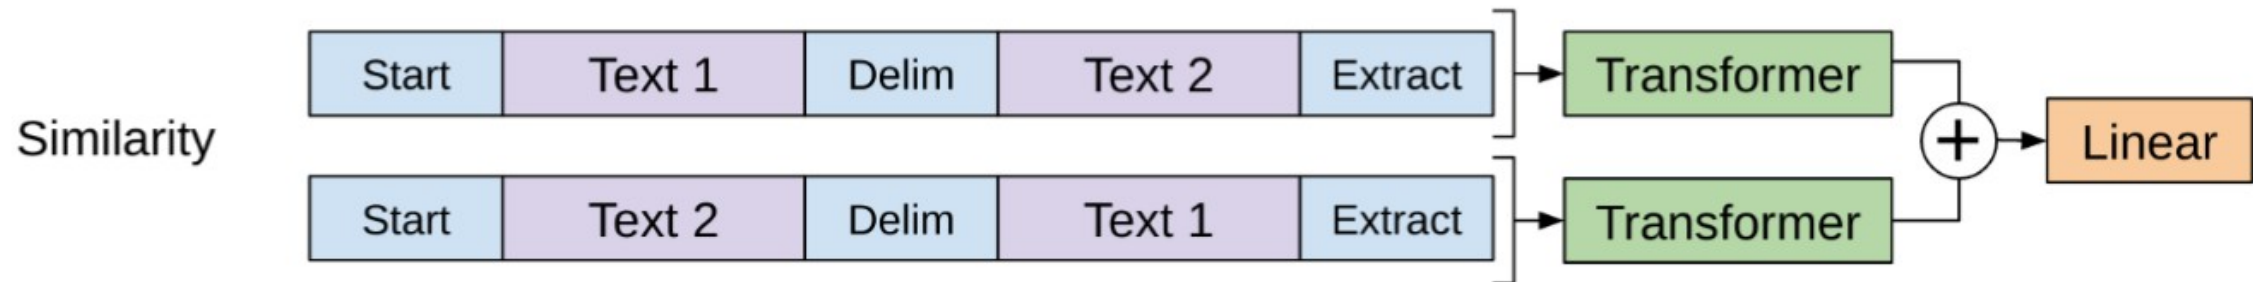

1 2 3 4 5 6 7 8 ... 512
[CLS] Let's stick to [MASK] in this skit

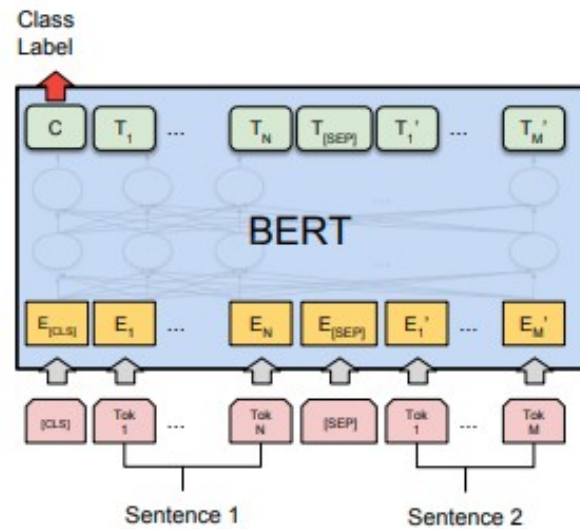
Input

[CLS] Let's stick to improvisation in this skit

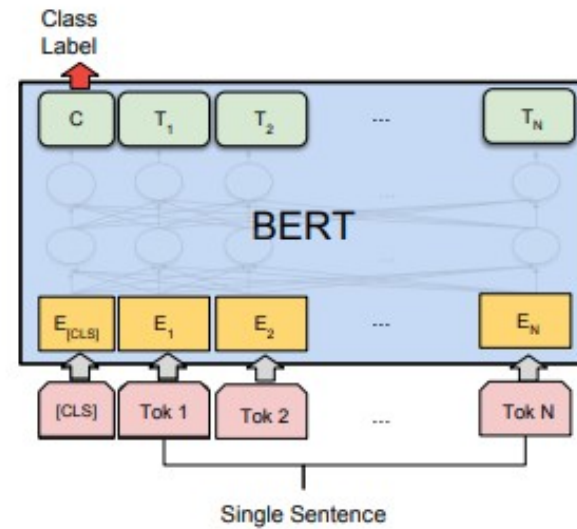
Predict likelihood
that sentence B
belongs after
sentence A



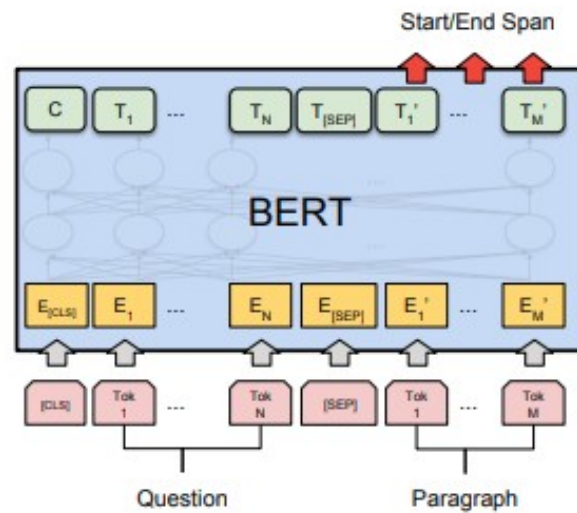




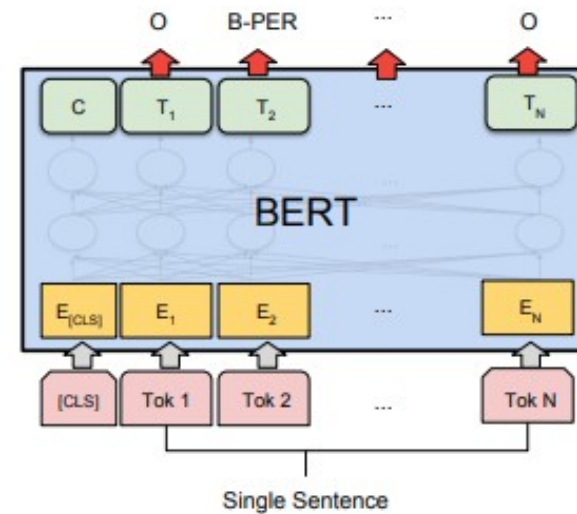
(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG



(b) Single Sentence Classification Tasks:
SST-2, CoLA



(c) Question Answering Tasks:
SQuAD v1.1

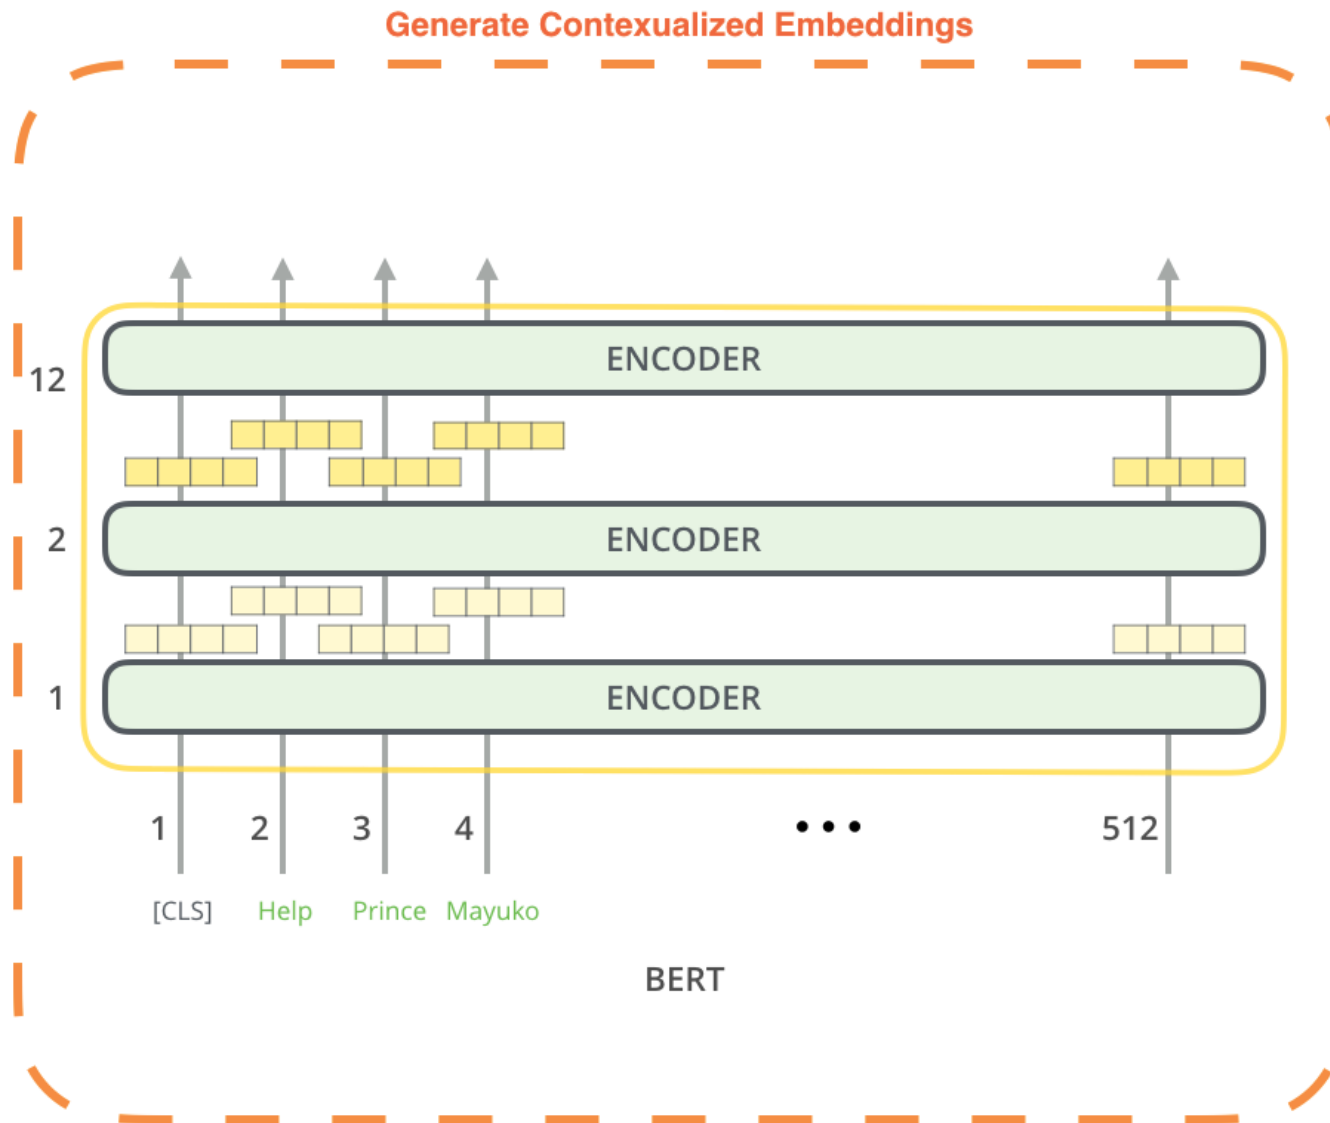


(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER


Training Bert

- Pre Training takes 14 days on a TPUv2 (500\$)
- Bert Large Models (24 Layers) can only be trained on TPUs
- Fine-tuning a model with 1GB of text takes several hours on a single GPU (1080 / 2080)

Bert as Embedding




The output of each encoder layer along each token's path can be used as a feature representing that token.



But which one should we use?

Bert as Embedding

What is the best contextualized embedding for “Help” in that context?
For named-entity recognition task CoNLL-2003 NER



Summary

- Pre-Trained on unlabeled data, fine tune on domain specific data
- Better language understanding
 - Context matters!
- Better results with less data
 - Latest hotness [MT-DNN Liu et al.](#)
 - No need for vast amounts of training data

Identify offensive language

using word vectors and
FastText



Identify offensive language

~~using word vectors and
FastText~~

using Bert



WHK Jobs

- Do cutting edge deep learning:
 - Dialog Systems aka Chat Bots
 - Speech Recognition
 - Speech Synthesis
 - Text Classification / Generation
 - Build Alexa and Mycroft skills